



LOAN DEFAULT PREDICTION AND INVESTMENT STRATEGIES IN ONLINE LENDING

Done By
Pritha Ghosh
Tejaswi Cherukuri
Anoop Gopalam

IDS - 572 | Data Mining for Business-Assignment 1
Loan Default prediction and Investment Strategies in Online Lending

IDS 572 - Assignment 1 - Part B
Pritha Ghosh, Anoop Gopalam, Tejaswi Cherukuri

5. Develop decision tree models to predict default.

(a) Split the data into training and validation sets. What proportions do you consider, why?

A: We consider a 70-30 split between the training and validation sets. We avoid a larger split because the proportion of charged-off loans is only 15% in comparison to the 85% of the fully paid loans.

(b) Train decision tree models (use both rpart, c50)

[If something looks too good, it may be due to leakage – make sure you address this]

What parameters do you experiment with, and what performance do you obtain (on training and validation sets)? Clearly tabulate your results and briefly describe your findings.

How do you evaluate performance – which measure do you consider, and why?

A: We have created the following decision tree models-

1. RPart tree: Complexity parameter = 0.00036, MinSplit=30, split = information
Variables used in tree construction:

```
Variables actually used in tree construction:
[1] avg_cur_bal          bc_open_to_buy
[3] bc_util              dti
[5] earliest_cr_line     emp_length
[7] grade               loan_amnt
[9] mo_sin_old_il_acct   mo_sin_old_rev_tl_op
[11] mo_sin_rcnt_rev_tl_op mo_sin_rcnt_tl
[13] mths_since_recent_inq num_actv_bc_tl
[15] num_il_tl            num_op_rev_tl
[17] num_rev_tl_bal_gt_0  pct_tl_nvr_dlq
[19] percent_bc_gt_75     purpose
[21] sub_grade            tax_liens
[23] tot_hi_cred_lim      total_bc_limit
[25] total_il_high_credit_limit total_rev_hi_lim
```

Performance:

Train Set : Accuracy=85.73%

```
> table(pred = predict(lcDT3, true=train$loan_status
```

 true

| pred | Fully Paid | Charged Off |
|-------------|------------|-------------|
| Fully Paid | 48324 | 7976 |
| Charged Off | 116 | 299 |

```
> mean(predTrn3 == lcdfTrn$loan_status)
```

```
[1] 0.8573217
```

Test Set: Accuracy = 85.08%

```
> table(pred = predict(lcDT3, true=test$loan_status)
```

 true

| pred | Fully Paid | Charged Off |
|-------------|------------|-------------|
| Fully Paid | 20618 | 3489 |
| Charged Off | 137 | 63 |

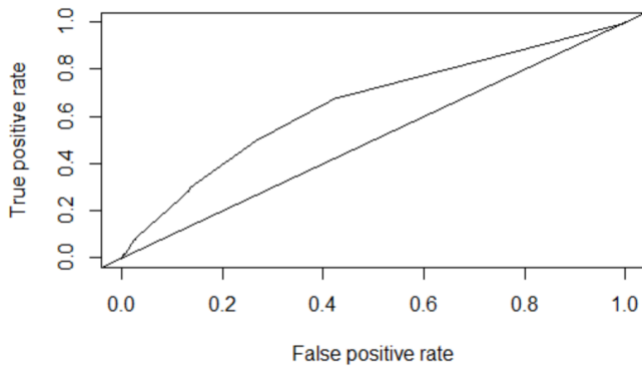
```
> mean(predict(lcDT3, lcdfTst, type='class') ==
```

```
[1] 0.8508249
```

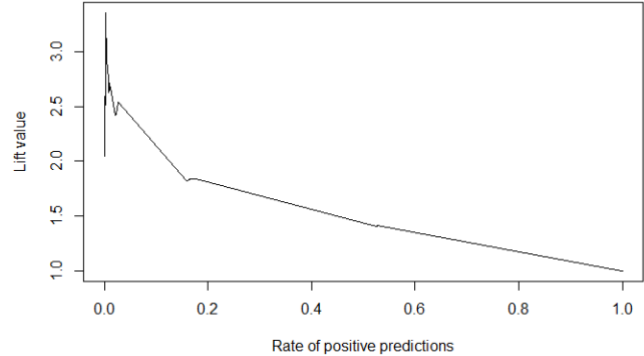
IDS - 572 | Data Mining for Business-Assignment 1

Loan Default prediction and Investment Strategies in Online Lending

AUC Curve: 0.647



Lift Curve:



2. RPart tree: Complexity parameter = 0.00036, MinSplit=30, split = gini

Variables used in tree construction:

Variables actually used in tree construction:

| | |
|----------------------------|----------------------------|
| [1] acc_open_past_24mths | annual_inc |
| [3] avg_cur_bal | bc_open_to_buy |
| [5] bc_util | dti |
| [7] earliest_cr_line | emp_length |
| [9] installment | int_rate |
| [11] loan_amnt | mo_sin_old_il_acct |
| [13] mo_sin_old_rev_tl_op | mo_sin_rcnt_rev_tl_op |
| [15] mort_acc | mths_since_recent_bc |
| [17] mths_since_recent_inq | num_accts_ever_120_pd |
| [19] num_actv_rev_tl | num_bc_sats |
| [21] num_il_tl | num_sats |
| [23] num_tl_op_past_12m | pct_tl_nvr_dlq |
| [25] purpose | sub_grade |
| [27] tot_hi_cred_lim | total_bal_ex_mort |
| [29] total_bc_limit | total_il_high_credit_limit |
| [31] total_rev_hi_lim | |

On Train Data: Accuracy = 85.92%

```
> table(pred = predTrn1, true=lcdfrn$loan_status)
      true
pred    Fully Paid Charged off
Fully Paid    48229      7803
Charged Off     179       504
> mean(predTrn1 == lcdfrn$loan_status)
[1] 0.8592612
```

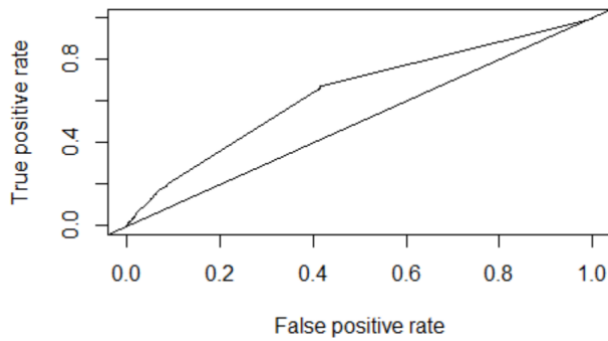
On Test Data: Accuracy = 85.16%

```
> table(pred = predTst, true=lcdftst$loan_status)
      true
pred    Fully Paid Charged off
Fully Paid    20599      3419
Charged Off     188       101
> mean(predict(lcdt1,lcdftst, type='class') == lcdftst$loan_status)
[1] 0.8516065
```

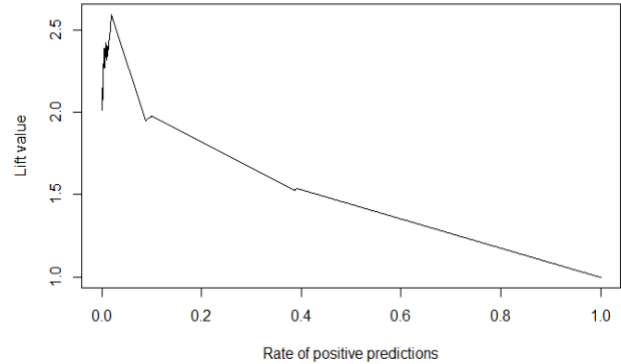
IDS - 572 | Data Mining for Business-Assignment 1

Loan Default prediction and Investment Strategies in Online Lending

AUC Curve:0.639



Lift Curve:



3. C50 tree:trials=50

Number of predictors = 49

```
>
> print(ctree)

Call:
C5.0.formula(formula = as.factor(trainset$loan_status)
~ ., data = trainset, method = "class", trials = 50)

Classification Tree
Number of samples: 56715
Number of predictors: 49

Number of boosting iterations: 50
Average tree size: 157

Non-standard options: attempt to group attributes
```

Performance:

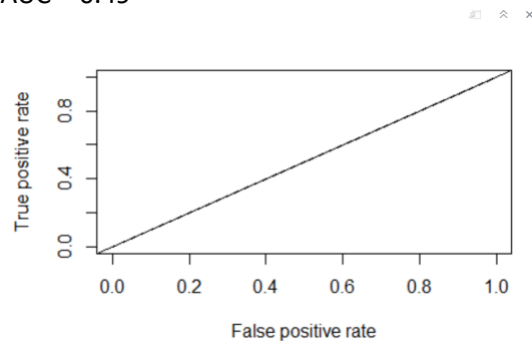
Train set : Accuracy = 86.54%

```
> table(pred = predTrain, true=trainset$loan.
      true
pred    Fully Paid Charged Off
Fully Paid    48467      7616
Charged Off     13      619
> mean(predTrain == trainset$loan_status)
[1] 0.8654853
> table(pred = predict(ctree, testset, type='c'
```

Test set: Accuracy:85.11%

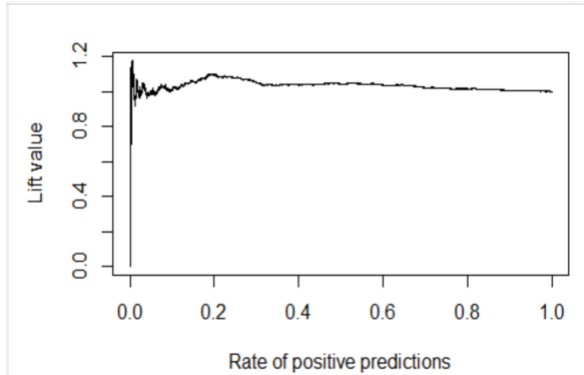
```
_status)
      true
pred    Fully Paid Charged Off
Fully Paid    20655      3557
Charged Off     60      35
> mean(predict(ctree, testset, type='c')
[1] 0.8511951
> |
```

AUC = 0.49



IDS - 572 | Data Mining for Business-Assignment 1
Loan Default prediction and Investment Strategies in Online Lending

Lift Curve:



To evaluate performance, we see the accuracy in predicting the response variable correctly. We also use the AUC and Lift curves as a test of performance.

(c) Identify the best tree model. Why do you consider it best? Describe this model – in terms of complexity (size). Examine variable importance. Briefly describe how variable importance is obtained in your best model.

A: For the first decision tree using rpart and information split, we got an accuracy of 85.73% on the training set and an accuracy of 85.08% on the test set. The AUC value that we got was 0.647. For the second decision tree using rpart and gini split, we got an accuracy of 85.92% on the training set and an accuracy of 85.16% on the test set. The AUC value that we got was 0.639. For the third decision tree using c50, we got an accuracy of 86.54% on the training set and an accuracy of 85.11% on the test set. The AUC value that we got was 0.49.

We reject the c50 tree because the AUC value is quite low.

Both the rpart trees perform considerably better. They both have AUC value of greater than 0.6 and perform well on both the train as well as the test sets.

If we would be forced to choose between one of them, we would choose the one with the gini split, because it performs slightly better, when we keep all the other parameter like cp and minsplitt the same. The one with information split would require more processing time.

The lift curve of that model is also used as a measure of the effectiveness of a predictive model, calculated as the ratio between the results obtained with and without the predictive model. The greater the area between the lift curve and the baseline, the better the model. The lift curve of the rpart tree with the gini split shows that the model is good.

This model uses 25 variables in the tree construction. This is far less than the number of variables used in the other trees, therefore we can be certain that our model will not cause overfit.

IDS - 572 | Data Mining for Business-Assignment 1
Loan Default prediction and Investment Strategies in Online Lending

The variable importance of this model is listed below :

| | |
|------------------|--------------|
| grade | 1083.7846556 |
| int_rate | 1026.3511686 |
| bc_open_to_buy | 237.8606938 |
| total_bc_limit | 212.5593145 |
| total_rev_hi_lim | 174.1314679 |
| sub_grade | 152.3769180 |
| annual_inc | 148.1297711 |

6. Develop a random forest model. (Note the 'ranger' library can give faster computations) What parameters do you experiment with, and does this affect performance? Describe the best model in terms of number of trees, performance, variable importance. Compare the random forest and best decision tree model from Q 4 above. Do you find the importance of variables to be different? Which model would you prefer, and why ?

In the first model, we have kept the num.trees = 50, importance = permutation.

We got an accuracy of 85% on the test data and a prediction error of 15%.

Confusion Matrix and Statistics

| | | |
|-------------|-------------|------------|
| | Charged Off | Fully Paid |
| Charged Off | 31 | 3504 |
| Fully Paid | 46 | 20726 |

Accuracy : 0.854
95% CI : (0.8494, 0.8584)
No Information Rate : 0.9968
P-Value [Acc > NIR] : 1

Kappa : 0.011

Mcnemar's Test P-Value : <2e-16

Sensitivity : 0.402597
Specificity : 0.855386
Pos Pred Value : 0.008769
Neg Pred Value : 0.997785
Prevalence : 0.003168
Detection Rate : 0.001275
Detection Prevalence : 0.145431
Balanced Accuracy : 0.628992

'Positive' class : Charged off

```
> rgModel2$prediction.error  
[1] 0.1510888
```

In the second model, we experiment with parameters num.trees=50 and importance = impurity.

| | | |
|-------------|-------------|------------|
| | Charged Off | Fully Paid |
| Charged Off | 37 | 3498 |
| Fully Paid | 43 | 20729 |

Accuracy : 0.8543
95% CI : (0.8498, 0.8587)
No Information Rate : 0.9967
P-Value [Acc > NIR] : 1

Kappa : 0.0141

Mcnemar's Test P-Value : <2e-16

Sensitivity : 0.462500
Specificity : 0.855616
Pos Pred Value : 0.010467
Neg Pred Value : 0.997930
Prevalence : 0.003291
Detection Rate : 0.001522
Detection Prevalence : 0.145431
Balanced Accuracy : 0.659058

```
> rgModel2$prediction.error  
[1] 0.1503659
```

IDS - 572 | Data Mining for Business-Assignment 1

Loan Default prediction and Investment Strategies in Online Lending

This model also gives us an accuracy of 85% on test data and a prediction error of 15%, but it does a slightly better job at predicting the “Charged Off” loans.

In the first model, the variable importance is as below -

```
> sort(rgModel1$variable.importance, decreasing = TRUE)
```

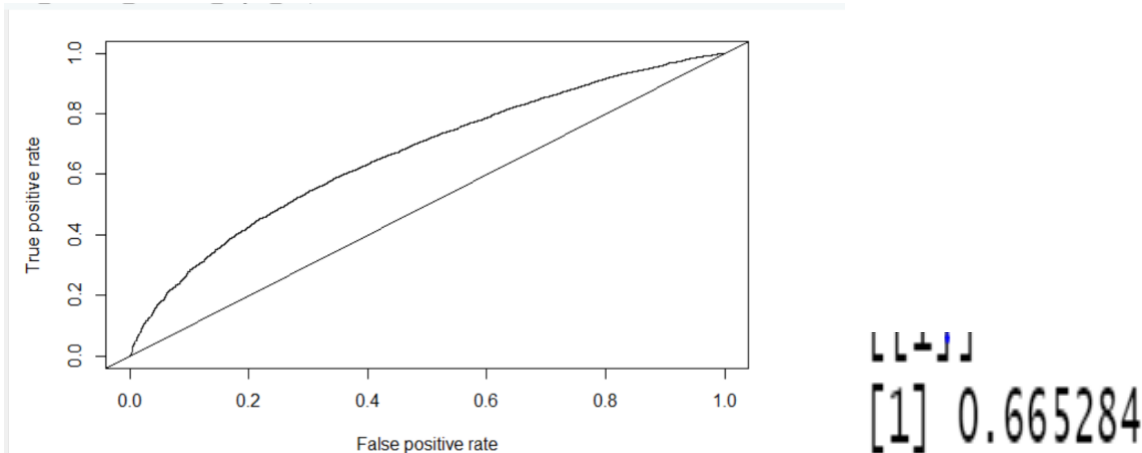
| | | |
|-----------------|------------------|--------------|
| tot_hi_cred_lim | total_bc_limit | avg_cur_bal |
| 1.753977e-02 | 1.359696e-02 | 1.245160e-02 |
| bc_open_to_buy | total_rev_hi_lim | installment |
| 1.236969e-02 | 1.100106e-02 | 8.754416e-03 |
| bc_util | annual_inc | loan_amnt |
| 8.704137e-03 | 8.166425e-03 | 7.976640e-03 |

In the second model, the variable importance is as below -

```
> sort(rgModel2$variable.importance*1000, decreasing = TRUE)
```

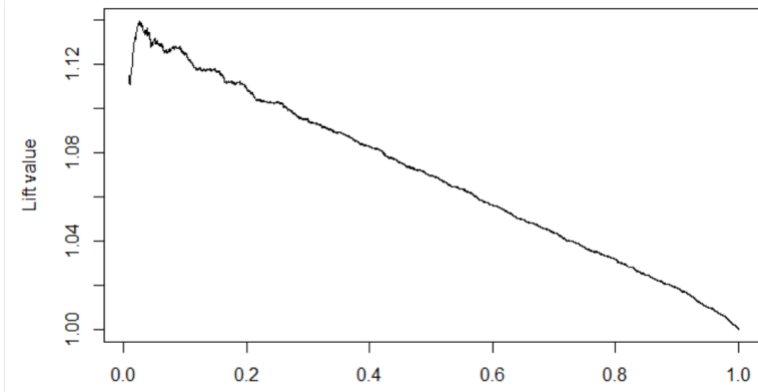
| | | |
|------------------|-----------------|----------------------|
| dti | int_rate | avg_cur_bal |
| 568674.519 | 516728.081 | 508467.502 |
| bc_open_to_buy | bc_util | mo_sin_old_rev_tl_op |
| 493121.742 | 491758.229 | 489483.699 |
| total_rev_hi_lim | tot_hi_cred_lim | mo_sin_old_il_acct |
| 480629.130 | 475113.646 | 472828.301 |

We choose the second model since it is comparatively better at predicting the Charged Off Loans
The AUC curve for the model is as below -



The AUC value is : 0.66.

Lift Curve:



IDS - 572 | Data Mining for Business-Assignment 1

Loan Default prediction and Investment Strategies in Online Lending

Since the model is not predicting the “Charged Off” cases correctly, we try oversampling the data.

```
> data_balanced_over <- ovun.sample(loan_status ~ ., data = lcdfTrn, method = "over", N = 95000)$data
> table(data_balanced_over$loan_status)
```

| | |
|-------------|------------|
| Charged Off | Fully Paid |
| 46603 | 48397 |

After oversampling, the confusion matrix on the test set is as below-

```

      Charged Off Fully Paid
Charged Off    117     3392
Fully Paid     258    20540

      Accuracy : 0.8498
      95% CI : (0.8453, 0.8543)
    No Information Rate : 0.9846
    P-Value [Acc > NIR] : 1

      Kappa : 0.0333

McNemar's Test P-Value : <2e-16

      Sensitivity : 0.312000
      Specificity : 0.858265
    Pos Pred Value : 0.033343
    Neg Pred Value : 0.987595
      Prevalence : 0.015428
    Detection Rate : 0.004813
    Detection Prevalence : 0.144362
    Balanced Accuracy : 0.585133

'Positive' Class : Charged Off

```

After balancing the dataset, the performance on predicting the “Charged Off” loans has improved by 2%.

Among the decision tree and random forest models, I would prefer the decision tree model with rpart() since it has the best overall accuracy and does a much better job at predicting the charged off loans.

Additional Model : XG Boost

Parameters: nrounds = 500, nfold=5, early_stopping_rounds = 10

```

[55] train-error:0.146663 train-auc:0.706721 eval-error:0.144362 eval-auc:0.681558
stopping. Best iteration:
[75] train-error:0.146663 train-auc:0.706721 eval-error:0.144362 eval-auc:0.681558

> xgb_1$ml$best_iteration

```

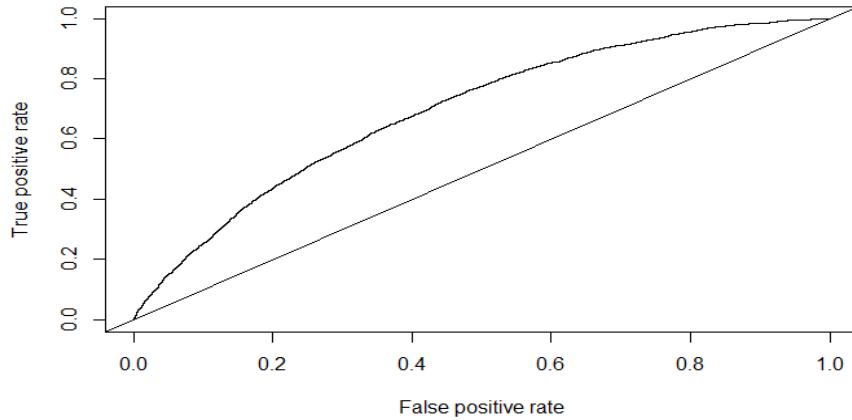
AUC value on test set: 0.68

Top 10 important variables -

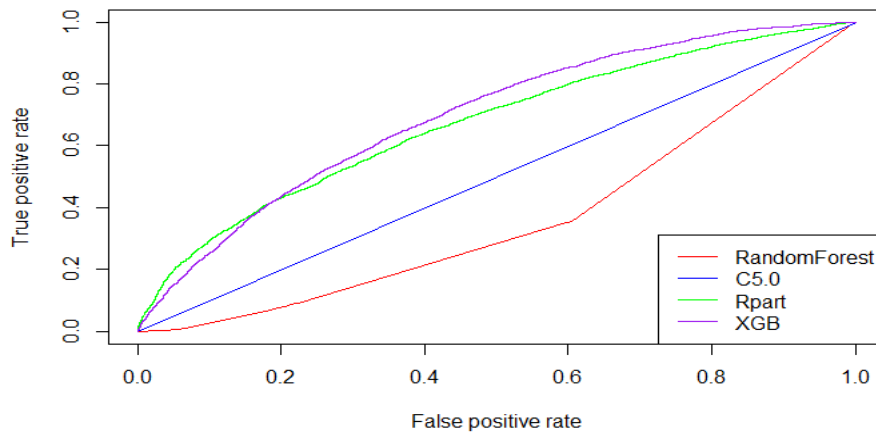
| | Feature | Gain | Cover | Frequency |
|----|----------------------|--------------|--------------|--------------|
| 1 | int_rate | 0.5465437685 | 1.976317e-01 | 0.1036525173 |
| 2 | dti | 0.0378911253 | 5.120722e-02 | 0.0523198421 |
| 3 | installment | 0.0370865508 | 5.768104e-02 | 0.0533070089 |
| 4 | acc_open_past_24mths | 0.0325252746 | 5.219352e-02 | 0.0365251728 |
| 5 | avg_cur_bal | 0.0325087844 | 4.197792e-02 | 0.0493583416 |
| 6 | tot_hi_cred_lim | 0.0316494806 | 7.171615e-02 | 0.0631786772 |
| 7 | annual_inc | 0.0199906575 | 4.195277e-02 | 0.0473840079 |
| 8 | bc_util | 0.0194640701 | 2.580485e-02 | 0.0434353406 |
| 9 | mo_sin_old_rev_tl_op | 0.0184014478 | 4.232077e-02 | 0.0444225074 |
| 10 | emp_length/a | 0.0145163083 | 2.980452e-02 | 0.0187561698 |

IDS - 572 | Data Mining for Business-Assignment 1
Loan Default prediction and Investment Strategies in Online Lending

AUC:



Consolidated ROC curves –



| Model | Parameters | Accuracy | AUC Value |
|----------------------------|--|-----------------------------|-----------|
| Rpart Decision Tree | Complexity parameter = 0.00036, MinSplit=30, split = gini | Train:85.92% Test:85.16% | 0.639 |
| C50 Decision Tree | trials=50 | Train:86.54 Test: 85.11% | 0.49 |
| Random Forest | num.trees = 50, importance = impurity | Train:87% Test:85.4% | 0.66 |
| XGBoost | nrounds = 500, nfold=5, early_stopping_rounds = 10 | Train:85.4% Test:85.6% | 0.68 |

IDS - 572 | Data Mining for Business-Assignment 1
Loan Default prediction and Investment Strategies in Online Lending

7. The average term for the Charged Off loans is 3 years, with an average interest rate of 13.8% and an average return rate of -11.69%.

The average term for the Fully Paid Off loans is 2.09 years, with an average interest rate of 11.56% and an average return rate of 8.03%.

| loan_status <chr> | avgInt <dbl> | avgRet <dbl> | avgTerm <dbl> |
|----------------------|-----------------|-----------------|------------------|
| Charged Off | 13.88228 | -11.697488 | 3.00000 |
| Fully Paid | 11.56970 | 8.031333 | 2.09523 |

Therefore, for Fully Paid loans, on an investment of \$100, there will be an average profit of \$8 per year.

Since the average term for the loan is 2.09 years, the total average return on investment would be ~\$16. To compare equivalently to charged off loans, which have an average term of 3 years, let us assume that for the remainder of the 1 year term remaining for the Fully paid loans, the return is at a risk-free rate of 2%.

Therefore, we can assume that for the total term of 3 years, one would make a profit of about \$18 on the Fully paid off loans.

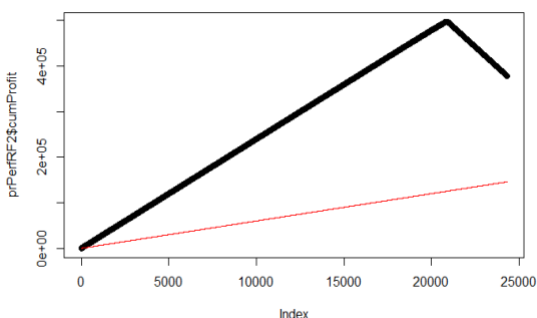
This is a conservative assumption, for ease of our calculations. In a real-life scenario, the average return is computed monthly and not annually, therefore, an investor is easily able to re-invest the profits gained on a loan into another loan to make even higher profits.

| decile <int> | count <int> | numDefaults <int> | avgActRet <dbl> | minRet <dbl> | maxRet <dbl> | avgTer <dbl> | totA <int> | totB <int> | totC <int> | totD <int> | totE <int> | totF <int> |
|-----------------|----------------|----------------------|--------------------|-----------------|-----------------|-----------------|---------------|---------------|---------------|---------------|---------------|---------------|
| 1 | 2431 | 0 | 8.3180140 | 0.0000000 | 31.17493 | 1.677289 | 819 | 770 | 595 | 220 | 26 | 1 |
| 2 | 2431 | 0 | 8.2328534 | 0.0000000 | 27.28570 | 1.716773 | 794 | 825 | 573 | 196 | 41 | 2 |
| 3 | 2431 | 0 | 8.4899999 | 3.0791241 | 26.79482 | 1.716853 | 804 | 763 | 595 | 234 | 33 | 2 |
| 4 | 2431 | 0 | 8.2877414 | 0.0000000 | 29.60077 | 1.875930 | 805 | 717 | 628 | 230 | 44 | 7 |
| 5 | 2431 | 1 | 7.8313673 | 0.8051210 | 31.09951 | 2.188751 | 699 | 725 | 683 | 255 | 62 | 7 |
| 6 | 2431 | 3 | 7.7750295 | -0.1458667 | 29.28361 | 2.254705 | 653 | 697 | 688 | 310 | 63 | 19 |
| 7 | 2431 | 4 | 7.7079562 | -0.3334643 | 29.01263 | 2.456243 | 541 | 741 | 737 | 296 | 96 | 16 |
| 8 | 2430 | 18 | 7.6009974 | -4.8463582 | 30.65945 | 2.565180 | 445 | 692 | 739 | 412 | 118 | 22 |
| 9 | 2430 | 1115 | -0.6306885 | -33.3333333 | 36.73610 | 2.812541 | 253 | 575 | 726 | 552 | 265 | 46 |
| 10 | 2430 | 2430 | -12.0736247 | -32.2695000 | 10.18747 | 3.000000 | 175 | 540 | 887 | 586 | 198 | 37 |

We first analyse the profits using the random forest method - This model is not ideal because it doesn't do a very good job at predicting the charged off loans correctly.

Score corresponding to maximum profit :

```
> which.max(prPerfRF2$cumProfit)
[1] 20841
```



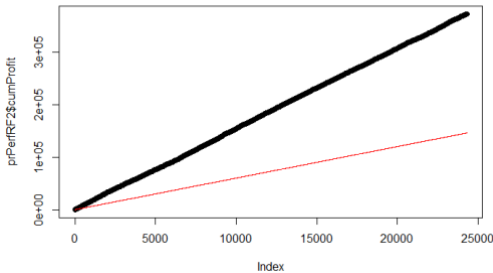
| | | | | |
|------|-----------|-------------|-----|--------|
| 7474 | 0.9960000 | Charged Off | -35 | 226082 |
|------|-----------|-------------|-----|--------|

IDS - 572 | Data Mining for Business-Assignment 1

Loan Default prediction and Investment Strategies in Online Lending

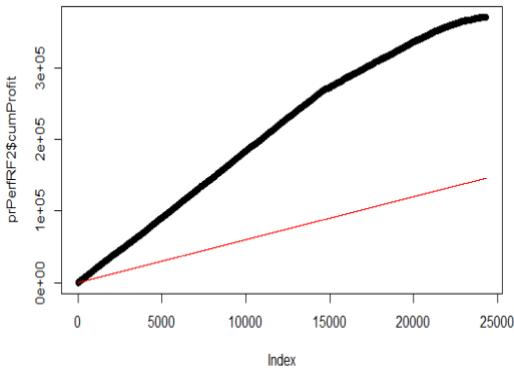
The threshold value is too high with the random forest model.

With C50 decision tree –



| scoreRF | status | profit | cumProfit |
|-----------|-------------|--------|-----------|
| 0.1454465 | Charged Off | -35 | 277 |

Next, we do the same analysis using the rpart decision tree - This is our best model as it does an equally good job of predicting both the Charged Off as well as the Fully Paid loans correctly.



| | scoreRF | status | profit | cumProfit |
|-------|-----------|-------------|--------|-----------|
| 1868 | 0.9090909 | Fully Paid | 24 | 24 |
| 3568 | 0.9090909 | Fully Paid | 24 | 48 |
| 5062 | 0.9090909 | Charged Off | -35 | 13 |
| 12692 | 0.9090909 | Fully Paid | 24 | 37 |
| 15693 | 0.9090909 | Fully Paid | 24 | 61 |

For this model, the threshold is 0.909. We observe this in the graph as well - up until about the 15000th prediction, the graph is rising steadily upwards. The rate decreases when the probability of prediction drops to less than ~90%, which is the optimal threshold.